


Speedup of GPU implementations



- Pytorch array programming
- CUDA-C kernel
- Pytorch Conv2d layer with stencil weights
- CUDA.jl stencil
- CUDA.jl array programming